

A Whole New Efficient Fuzzing Strategy for Stagefright

Porting and Optimizations

by Zinuo Han
at Ruxcon 2017

About Me

Security Researcher at Chengdu Security Response Center, Qihoo 360

Focused on Android vulnerability research and exploit development

ele7enxxh@gmail.com

<https://github.com/ele7enxxh>

Agenda

Introduction

Design

Implementation

Disclosure

Conclusion

Introduction

Basic information about this talk

What is Stagefright

A logical algorithm framework library for parsing multimedia on Android

- Written in C++

Support a wide range of audio and video formats

- Including MP3, MP4, MKV, MPEG2, MPEG4, and many more

Lots of vulnerabilities have already been found in Stagefright

Why Stagefright again

Stagefright vulnerabilities keep attractive

- Continuously patched in every Android security update since August 2015

Most Stagefright vulnerabilities assessed as Critical

- This means higher bounty

Found vulnerabilities in multiple Stagefright components

- Libstagefright library, especially MPEG4Extractor.cpp
- OMX
- SW codecs(Most in 2017)

Why Stagefright again

Stagefright vulnerabilities keep attractive

- Continuously patched in every Android security update since August 2015

Most Stagefright vulnerabilities assessed as Critical

- This means higher bounty

Found vulnerabilities in multiple Stagefright components

- Libstagefright, especially MPEGExtractor.cpp
- OMX
- SW codecs(Most in 2017)

Still 0days

Related Work

| Research | Focused on | Coverage -Guide | ASAN | Fast | Crash Tolerant |
|--|----------------|--------------------|------|------|-------------------|
| Fuzzing the Media Framework in Android(MFFA) by Alexandru Blanda at ELC 2015 | libstagefright | ✗ | ✗ | - | ✓ |
| | | | | | |
| | | | | | |

Related Work

| Research | Focused on | Coverage -Guide | ASAN | Fast | Crash Tolerant |
|--|----------------|--------------------|------|------|-------------------|
| Fuzzing the Media Framework in Android(MFFA) by Alexandru Blanda at ELC 2015 | libstagefright | ✗ | ✗ | - | ✓ |
| Stagefright: Scary Code in the Heart of Android by Joshua Drake at Blackhat USA 2015 | MPEG4Extractor | ✓ | ✓ | + | ✓ |
| | | | | | |

Related Work

| Research | Focused on | Coverage -Guide | ASAN | Fast | Crash Tolerant |
|--|----------------|--------------------|------|------|-------------------|
| Fuzzing the Media Framework in Android(MFFA) by Alexandru Blanda at ELC 2015 | libstagefright | ✗ | ✗ | - | ✓ |
| Stagefright: Scary Code in the Heart of Android by Joshua Drake at Blackhat USA 2015 | libstagefright | ✓ | ✓ | + | ✓ |
| Fuzzing Android OMX by MingjianZhou and ChiachihWu at HITCON 2016 | OMX | ? | ? | + | ✓ |

About this talk

What will be talked about next

- How to design and implement a efficient fuzzing strategy for Stagefright
- What vulnerabilities did I find by the above method
- Conclusion of this talk

What will not be talked about next

- The root cause of the vulnerabilities
- Vulnerabilites exploitation

Design

Goal & Architecture overview

Goal

More targeted

- Mainly focus on SW codecs

Goal

More targeted

- Mainly focus on SW codecs

More faster

- Run Stagefright on desktop Linux
- Optimize Stagefright workflow

Goal

More targeted

- Mainly focus on SW codecs

More faster

- Run Stagefright on desktop Linux
- Optimize Stagefright workflow

More technical

- Coverage-Guided fuzzer(American Fuzzy Lop)
- AddressSanitizer

Goal

More targeted

- Mainly focus on SW codecs

More faster

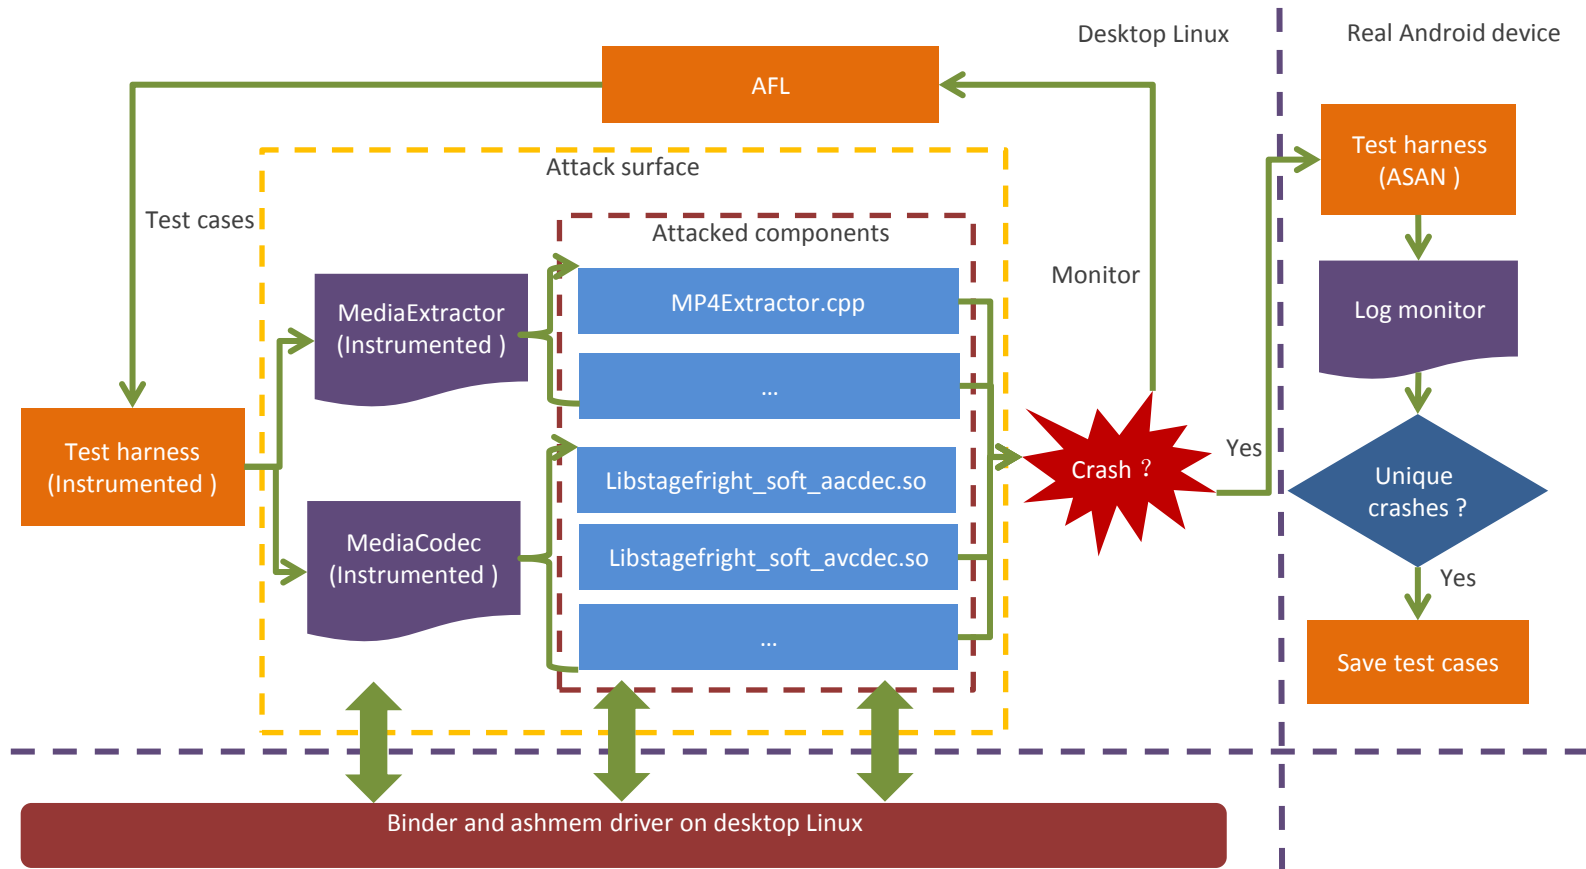
- Run Stagefright on desktop Linux
- Optimize Stagefright workflow

More technical

- Coverage-Guided fuzzer(American Fuzzy Lop)
- AddressSanitizer

Find 0days more easily

Architecture overview



Implementation

Details of the fuzzing strategy

Steps in the fuzzing strategy

Find attack surface

Porting

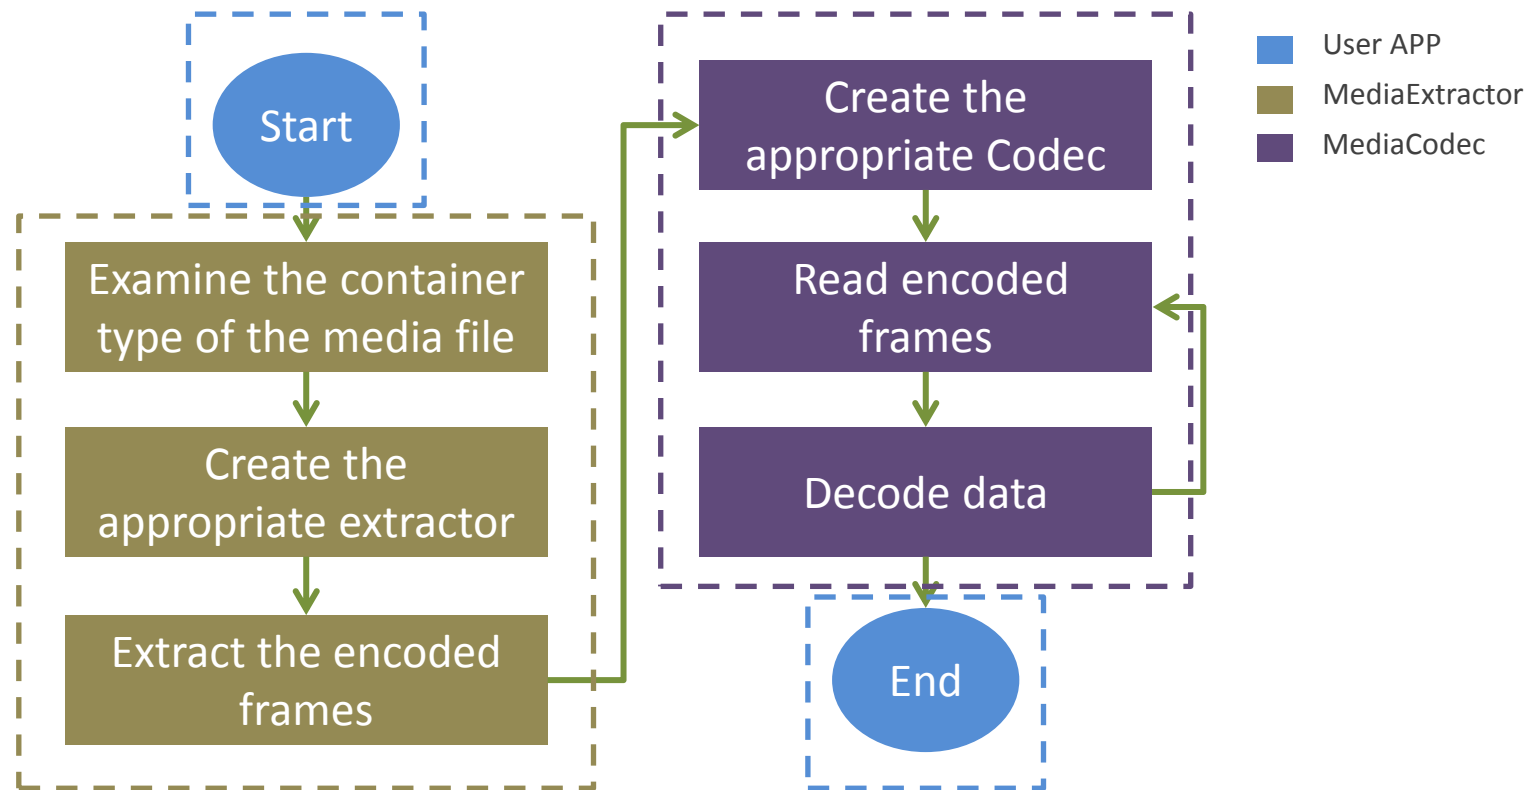
Optimizations

Get more powerful test cases

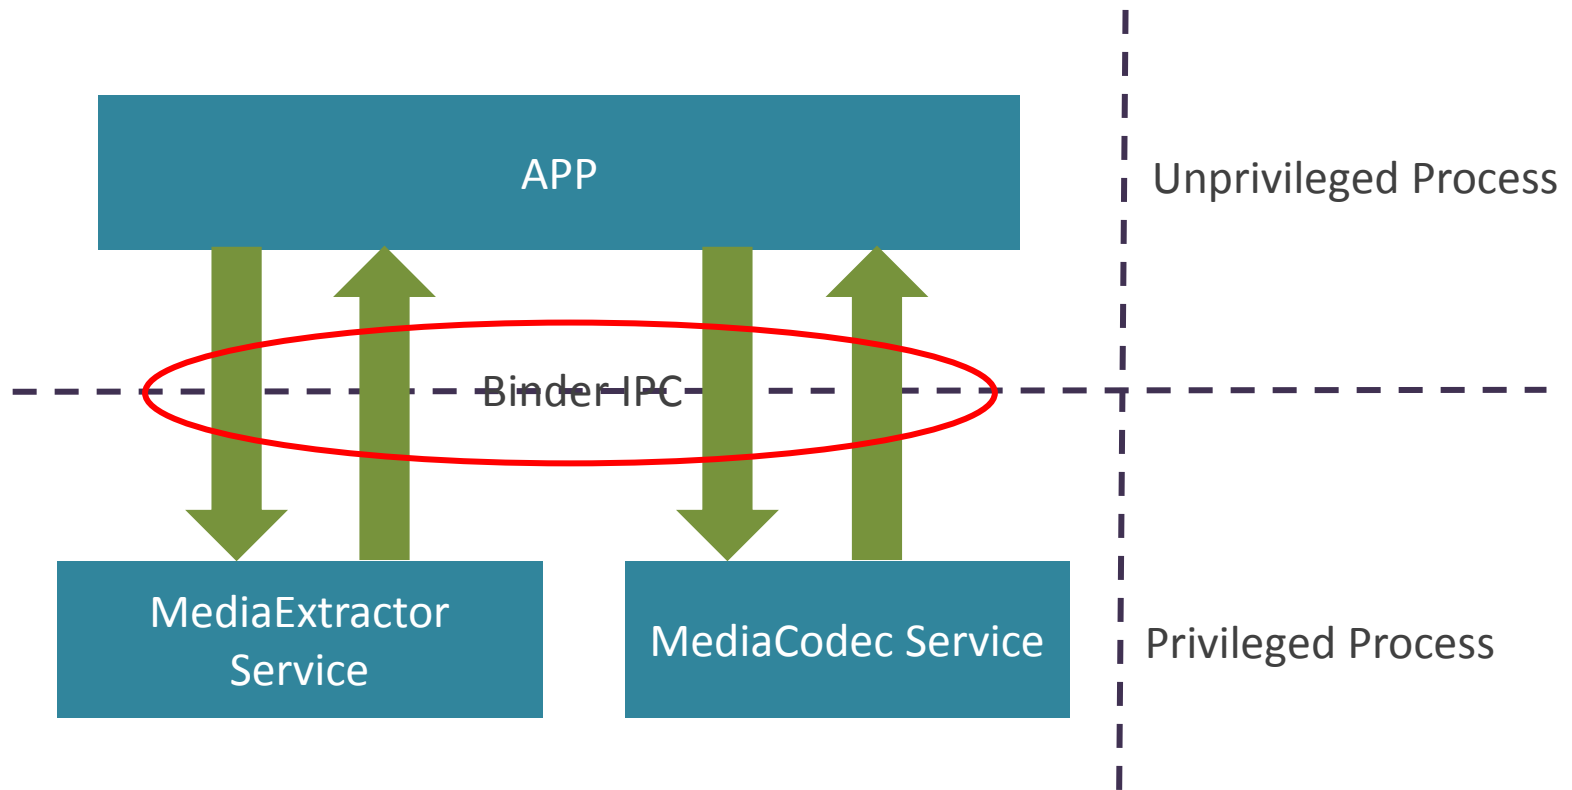
Fuzzing mp4 container

Recognize unique crashes

Find attack surface – How is audio and video played



Find attack surface – How is audio and video played



Find Attack Surface – Attacked components

| Container Type | Module | Codec Type | ASAN |
|----------------|-----------------------|----------------|----------------------------------|
| MP4 | MPEG4Extractor.cpp | AAC | libstagefright_soft_aacdec.so |
| MP3 | MP3Extractor.cpp | AMRNB AMRWB | libstagefright_soft_amrdec.so |
| AMRNB AMRWB | AMRExtractor.cpp | H264 | libstagefright_soft_avcdec.so |
| FLAC | FLACExtractor.cpp | HEVC | libstagefright_soft_hevcdec.so |
| WAV | WAVExtractor.cpp | G711 | libstagefright_soft_g711dec.so |
| OGG | OGGExtractor.cpp | MPEG2 | libstagefright_soft_mpeg2dec.so |
| MKV | MatroskaExtractor.cpp | H263 MPEG4 | libstagefright_soft_mpeg4dec.so |
| MPEG2TS | MPEG2TSExtractor.cpp | MP3 | libstagefright_soft_mp3dec.so |
| WVM | WVMExtractor.cpp | VORIBS | libstagefright_soft_vorbisdec.so |
| AAC | AACExtractor.cpp | OPUS | libstagefright_soft_opusdec.so |
| MPEG2PS | MPEG2PSExtractor.cpp | VP8 VP9 | libstagefright_soft_vpxdec.so |
| MIDI | MIDIExtractor.cpp | GSM | libstagefright_soft_gsmdec.so |

Find Attack Surface – Summary

Two potential attacked processes with privilege: **mediaextracor** and **mediacodec**

- Mediaextractor is used to extract audio and video frames
- Mediacodec is used to encode and decode audio and video frames

Full controllable input data

Complex input formats

- It means the possibility of more vulnerabilities

Easy to trigger

- No special permissions required

Porting – What and why

Port Stagefright to x86

- Android device most likely uses the ARM architecture, but the x86 family of processors is used in most desktop

Port binder and ashmem driver to linux

- Stagefright works with binder and ashmem driver, which are not enabled in desktop Linux as default

Port AFL to Android toolchains

- The shmat() function is used in the afl-llvm-rt.o.c, however Android toolchains can't recognize it

Setup running environment

Porting – What and why

Port Stagefright to x86

- Android device most likely uses the ARM architecture, but the x86 family of processors is used in most desktop

Port binder and ashmem driver to linux

- Stagefright works with binder and ashmem driver, which are not enabled in desktop Linux as default

Port AFL to Android toolchains

- The shmat() function is used in the afl-llvm-rt.o.c, however Android toolchains can't recognize it

Setup running environment

Make it efficient

Porting – Port Stagefright to x86

Download AOSP code from: <https://android.googlesource.com/>

Build Android for x86

```
$ cd aosp  
$ source build/envsetup.sh  
$ lunch aosp_x86-eng  
$ make -j64
```

Build Stagefright to x86

```
$ cd framework/av/cmds/stagefright  
$ make -j64
```

Porting – Port binder and ashmem driver to Linux

Download the latest version of the Linux kernel from:
<http://www.kernel.org>

Enable the binder and ashmem driver

- e.g. (Linux kernel version 4.8.17)

```
CONFIG_ANDROID=y && CONFIG_ANDROID_BINDER_IPC=y && CONFIG_ASHMEM=y
```

Add new udev rules to set correct permissions

```
$ echo -e "KERNEL==\"binder\",  
MODE=\"0666\"\\nKERNEL==\"ashmem\", MODE=\"0666\"\" |  
sudo tee /etc/udev/rules.d/android.rules
```

Porting – Port AFL to Android toolchains

Use `syscall()` function instead of `shmat()` function in `afl-llvm-rt.o.c`

```
- shmat(shm_id, NULL, 0);  
+ syscall(SYS_ipc, IPCOP_shmat, id, flag, &addr, addr);
```

Force the compile wrapper(`afl-clang-fast`) to instrument the shared libraries

```
- if (!strcmp(cur, "-shared")) maybe_linking = 0;
```

Cross-compile `afl-llvm-rt.o` which should be linked into the target Android binary

Porting – Setup running environment

Create a soft link

```
$ ln -s out/system /system
```

Copy configuration files

```
$ cp out/system/etc/media_codecs_google_audio.xml /etc
$ cp out/system/etc/media_codecs_google_telephony.xml /etc
$ cp out/system/etc/media_codecs_google_video.xml /etc
$ cp out/system/etc/media_codecs.xml /etc
```

Startup dependency services: e.g. mediaextractor, mediacodec...

Porting – It works

Running 5x times faster

```
sailfish:/data/local/tmp # time stagefright -s Disco.240p.mp4
thumbnailTime: 0 us (0.00 secs)
AVC video profile 66 and level 13
format changed.
.....$
avg. 63.91 fps
avg. time to decode one buffer 12987.23 usecs
decoded a total of 304 frame(s).
0m05.08s real    0m01.18s user    0m01.10s system
```

```
ele7enxxh@360:~$ lsb_release -a | grep Description
Description:    Ubuntu Yakkety Yak (development branch)
ele7enxxh@360:~$ time /system/bin/stagefright -s Disco.240p.mp4
thumbnailTime: 0 us (0.00 secs)
AVC video profile 66 and level 13
format changed.
.....$
avg. 180.82 fps
avg. time to decode one buffer 5473.04 usecs
decoded a total of 304 frame(s).

real    0m1.724s
user    0m0.352s
sys     0m0.032s
```

Porting – what else

Make it more efficient

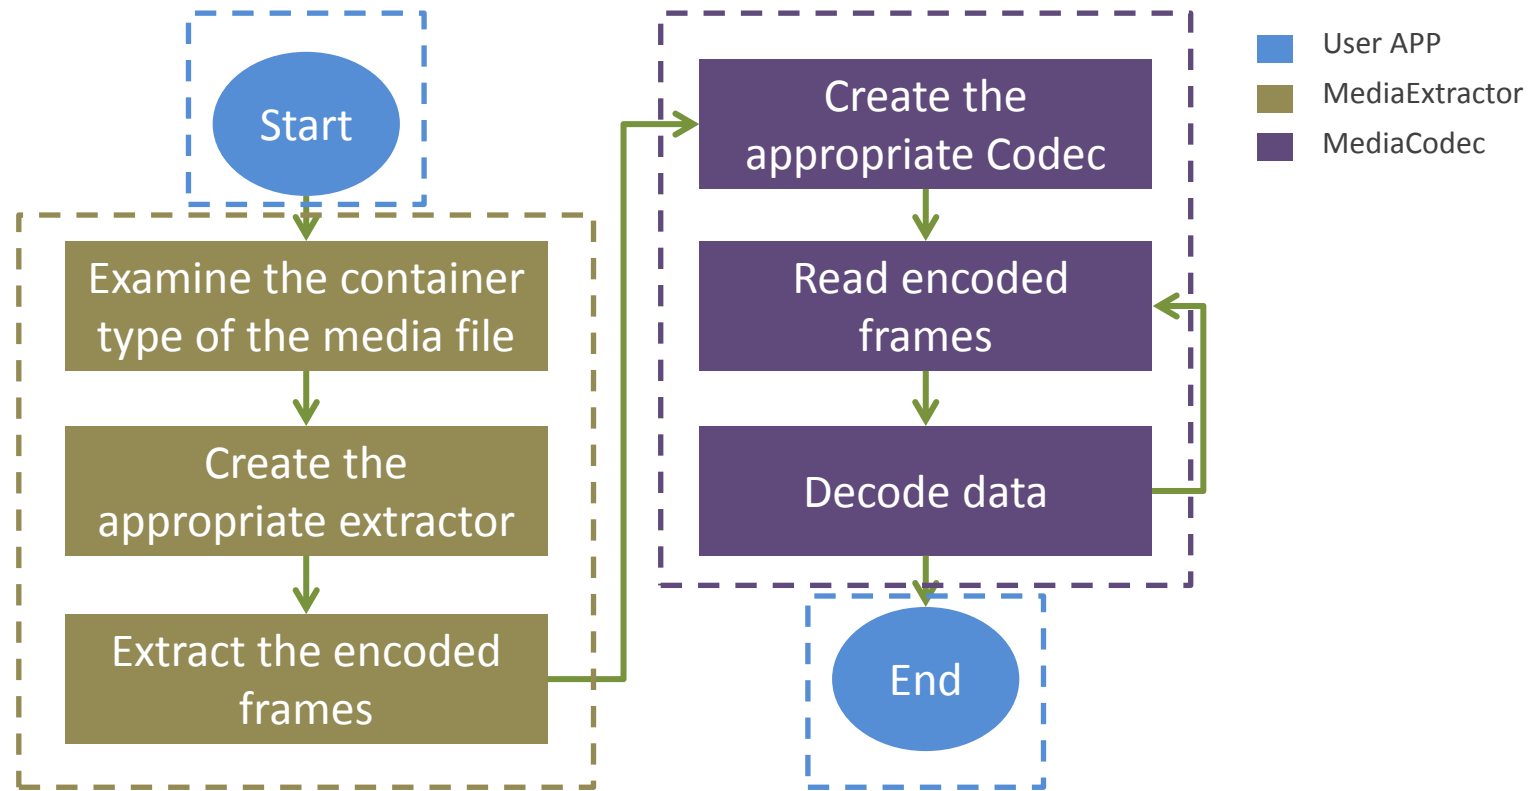
- Make extractor and codec work independently
- Bypass the media type sniffing mechanism
- Decode only one encoded frame

Make it more AFL-friendly

- Running as a single process

Let's start optimizing

Optimizations – Take a step back to Stagefright workflow



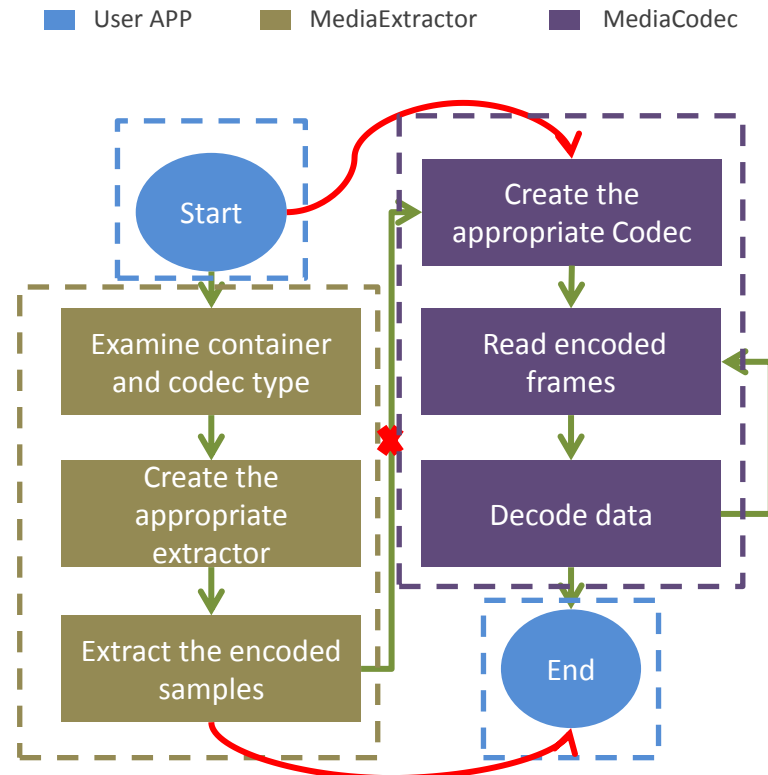
Optimizations – Make extractor and codec work independently

Why

- When fuzzing target is `codec(extractor)`, `extractor(codec)` will waste CPU time
- When fuzzing target is codec, unnecessary check logic in extractor could stop the decoding process ahead of time

How

- Fuzz only one target at a time



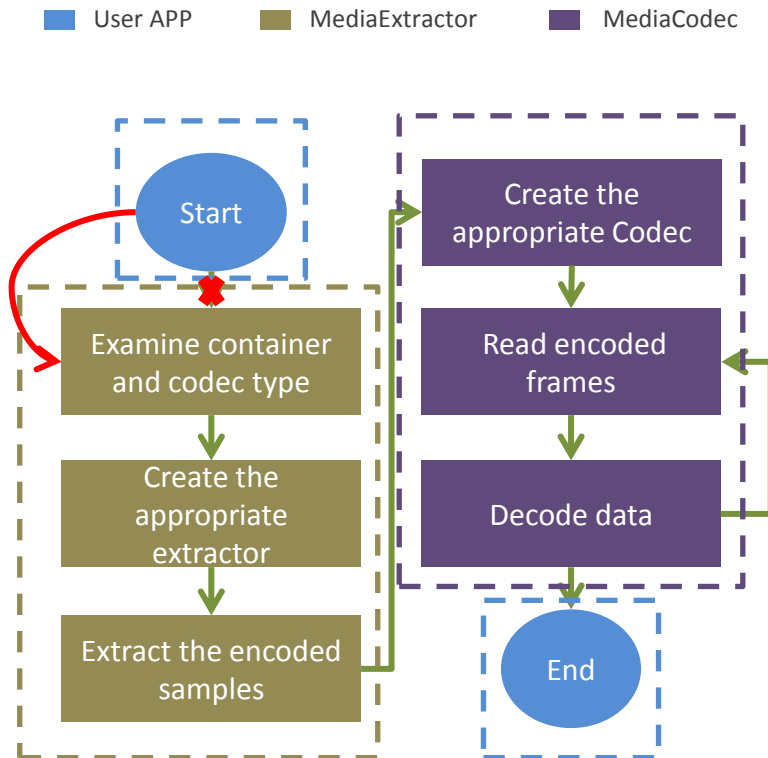
Optimizations – Bypass the sniffing mechanism

Why

- Fuzz one specify media type at a time is more efficient, the sniffing is not necessary for this job
- Cause meaningless mutations: there is chance that one media type could be turned into other media type by AFL

How

- Specify the container and codec type



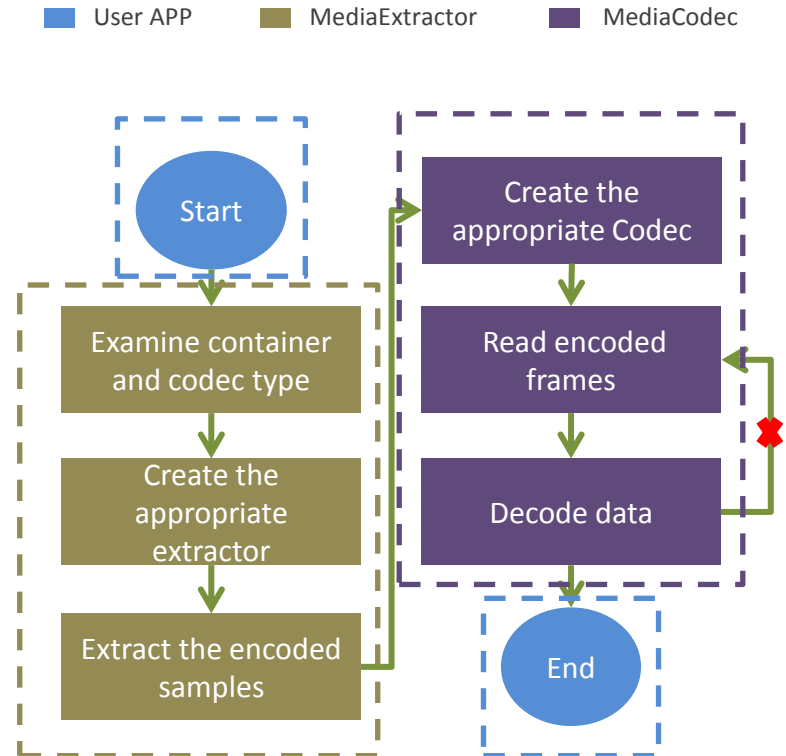
Optimizations – Decode only one frame

Why

- Typically, a standard media file contains multiple frames
- However, most vulnerabilities have been triggered when decoding the first frames – Just in my experience

How

- Break the decoding loop
- ```
if (numFrames == 1) break;
```



# Optimizations – Running as a single process

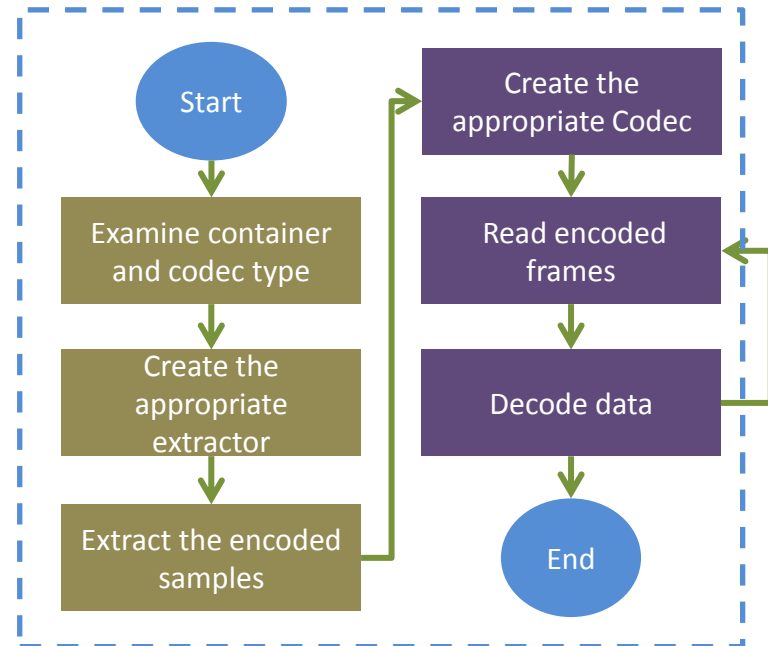
## Why

- Multi-process communication brings extra overhead
- AFL-unfriendliness

## How

- Create services in local

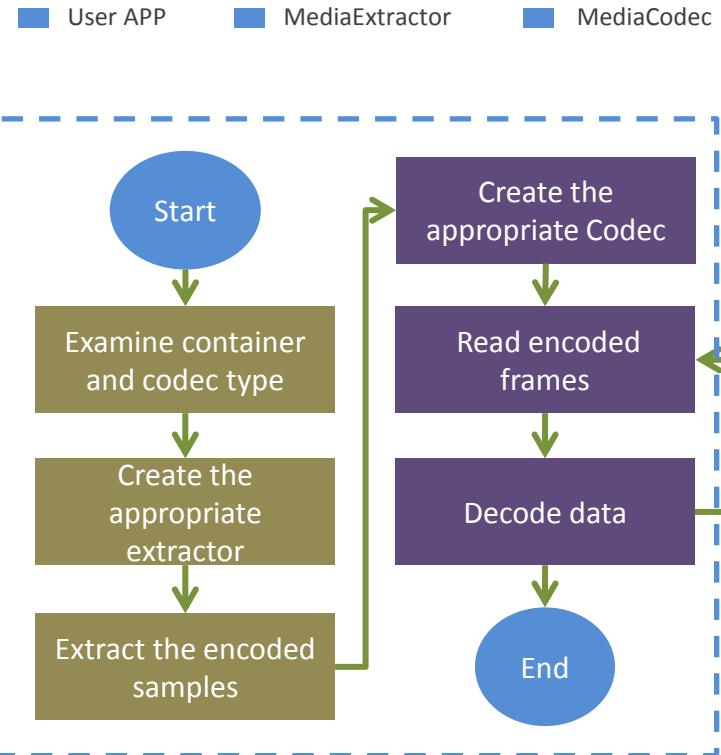
■ User APP   ■ MediaExtractor   ■ MediaCodec



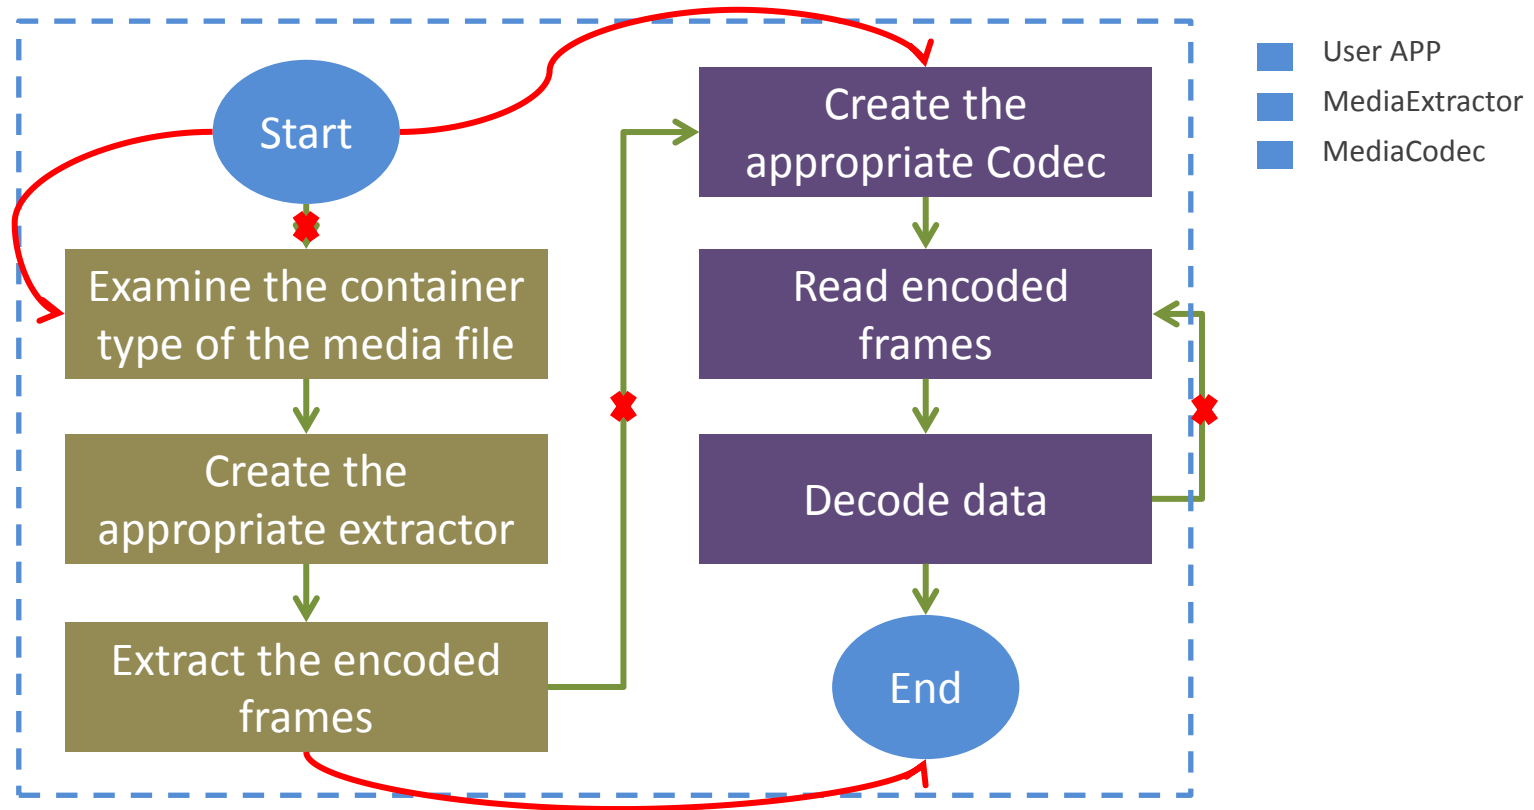
# Optimizations – Running as a single process

```
// Create MediaCodecList object instance in local
sp<IMediaCodecList> MediaCodecList::getLocalInstance()
{
 sCodecList = new MediaCodecList;
 return sCodecList;
}

sp<IMediaCodecList> MediaCodecList::getInstance() {
 // Get the remote interface of the mediaserver via
 binder IPC
 sp<IBinder> binder = defaultServiceManager()-
 >getService(String16("media.player"));
 sp<IMediaPlayerService> service =
 interface_cast<IMediaPlayerService>(binder);
 if (service.get() != NULL) {
 sRemoteList = service->getCodecList();
 if (sRemoteList == NULL) {
 // if failed to get remote list, create local
 list
 sRemoteList = getLocalInstance();
 }
 }
 return sRemoteList;
}
```



# Optimizations – The new Stagefright work flow



# Optimizations – It works faster

---

Two test harness for extractor and codec

20x+ performance gains

```
le7enxxh@360:~$ lsb_release -a | grep Description
Description: Ubuntu Yakkety Yak (development branch)
le7enxxh@360:~$ time /system/bin/stagefright -s Disco.240p.mp4
thumbnailTime: 0 us (0.00 secs)
AVC video profile 66 and level 13
format changed.
.....$
avg. 180.82 fps
avg. time to decode one buffer 5473.04 usecs
decoded a total of 304 frame(s).

real 0m1.724s
user 0m0.352s
sys 0m0.032s
```

```
le7enxxh@360:~$ lsb_release -a | grep Description
Description: Ubuntu Yakkety Yak (development branch)
le7enxxh@360:~$ time /system/bin/extractorfuzz video/mp4 Disco.240p.mp4
Create MPFG4Extracotr in local
real 0m0.031s
user 0m0.028s
sys 0m0.000s
le7enxxh@360:~$ time /system/bin/codecfuzz video/avc Disco.240p.mp4.avc
the code type is video/avc$

real 0m0.042s
user 0m0.036s
sys 0m0.004s
```

# Get more powerful test cases

## – Where to get test cases

---

The AOSP repository contains a large number of test cases

Search in Google

e.g. `-inurl:htm -inurl:html intitle:"index of" .mp4`

File format conversion tools

- ffmpeg

Open source project

- <https://github.com/MozillaSecurity/fuzzdata>
- <https://github.com/stribika/afl-fuzz/tree/master/testcases>



# Get more powerful test cases

## – How to get encoded frames

---

Dump the data when decoding a media file

Refer to the following code

```
//
framework/av/media/libstagefright/codecs/avcdec/SoftAV
CDec.cpp
void SoftAVC::onQueueFilled(OMX_U32 portIndex) {
 ...
 // If input dump is enabled, then write to file
 // pv_stream_buffer points to encoded frames
 // u4_num_Bytes is encoded frames length
 DUMP_TO_FILE(mInFile, s_dec_ip.pv_stream_buffer,
s_dec_ip.u4_num_Bytes);
 ...
}
```

# Get more powerful test cases

- Keep it small (< 1kb)

---

Keep only the video or audio data that actually want to fuzz

```
$ ffmpeg -i input_file -vcodec copy -an output_file_video
$ ffmpeg -i input_file -acodec copy -vn output_file_audio
```

Keep only one frame data

```
$ ffmpeg -i input_file -codec copy -frames 1 output_file
```

Use afl-tmin tool

# Get more powerful test cases

- Keep it small (< 1kb)

---

Keep only the video or audio data that you actually want to fuzz

```
$ ffmpeg -i input_file -vcodec copy -an output_file_video
$ ffmpeg -i input_file -acodec copy -vn output_file_audio
```

Keep only one frame data

```
$ ffmpeg -i input_file -codec copy -frames 1 output_file
```

Use afl-tmin tool

Remove useless data, save CPU time

# Fuzzing mp4 container – Test harness

---

```
// extractorfuzz.cpp
int main(int argc, char **argv) {
 if (argc != 3) return -1;
 ProcessState::self()->startThreadPool();
 // Running in persistent mode
 while (__AFL_LOOP(1000)) {
 // argv[2] is the input file path
 sp<FileSource> fileSource = new
FileSource(argv[2]);
 // argv[1] is the container type of the input file
 sp<IMediaExtractor> extractor =
MediaExtractor::Create(fileSource, argv[1]);
 }
 return 0;
}
```

# Fuzzing mp4 container – Instrument

---

Add the following to Android.mk in test harness and libstagefright

```
LOCAL_CLANG := true
LOCAL_CFLAGS += -fno-omit-frame-pointer
-O2
export AFL_PATH :=
/usr/local/lib/afl/arm
export AFL_CC := /usr/local/bin/clang
LOCAL_CC := afl-clang-fast
export AFL_CXX :=
/usr/local/bin/clang++
LOCAL_CXX := afl-clang-fast++
```

**Build**

```
$ mm -j16
```

# Fuzzing mp4 container – Get test cases

---

Get original test cases from here

- `cts/tests/tests/media/res/raw/a_4_aac.mp4`
- `cts/tests/tests/media/res/raw/swirl_128x128_h264.mp4`

Keep only one frame

```
ffmpeg -i a_4_aac.mp4 -codec copy -frames 1
a_4_aac_1frame.mp4
```

```
ffmpeg -i swirl_128x128_h264.mp4 -codec copy -frames 1
swirl_128x128_h264_1frame.mp4
```

# Fuzzing mp4 container – Lunch afl-fuzz loop

---

## Fuzzing in distributed mode

- Make full use of CPU performance

## Increase the -m and -t limits

- The decoding process requires more memory and time

```
$ afl-fuzz -M fuzz0 -m 1024 -t 1000 -i in -o out --
extractorfuzz video/mp4 @@
```

```
$ afl-fuzz -S fuzz1 -m 1024 -t 1000 -i in -o out --
extractorfuzz video/mp4 @@
```

...

# Fuzzing mp4 container – Running screen

## Faster

- Exec speed > 10k/sec

## Efficient

- Find 1400+ paths in 9 seconds

```
american fuzzy lop 2.51b (fuzz1)

- process timing -
 run time : 0 days, 0 hrs, 0 min, 9 sec
 last new path : 0 days, 0 hrs, 0 min, 0 sec
 last uniq crash : none seen yet
 last uniq hang : none seen yet
- cycle progress -
 now processing : 0 (0.00%)
 paths timed out : 0 (0.00%)
- stage progress -
 now trying : calibration
 stage execs : 0/8 (0.00%)
 total execs : 27.2k
 exec speed : 11.8k/sec
- fuzzing strategy yields -
 bit flips : n/a, n/a, n/a
 byte flips : n/a, n/a, n/a
 arithmetics : n/a, n/a, n/a
 known ints : n/a, n/a, n/a
 dictionary : n/a, n/a, n/a
 havoc : 1203/8192, 236/6496
 trim : 0.00%/460, n/a

- overall results -
 cycles done : 0
 total paths : 1479
 uniq crashes : 0
 uniq hangs : 0
- map coverage -
 map density : 1.02% / 4.78%
 count coverage : 1.42 bits/tuple
- findings in depth -
 favored paths : 2 (0.14%)
 new edges on : 1270 (85.87%)
 total crashes : 0 (0 unique)
 total tmouts : 0 (0 unique)
- path geometry -
 levels : 2
 pending : 1478
 pend fav : 1
 own finds : 1439
 imported : 37
 stability : 99.04%

[cpu001: 31%]
```



# Recognize unique crashes – Why

---

The AFL recorded crashes may be non-reproducible

- Different code for different processor architectures

Some crashes not interesting

- Assertion

The unique crashes recorded by AFL are not always unique

- AFL's uniqueness was determined based on tuple instrumentation that is too strict

Not crash  $\neq$  Not vulnerability

- Some vulnerabilities(e.g. use-after-free) don't cause crashes

# Recognize unique crashes – How

---

Push the unique crashes and the generated corpus to a real Android device with the latest security updates

Examine all corpus with ASAN-enabled again

- Build AOSP with ASAN:

```
$ make -j42
$ SANITIZE_TARGET=address make -j42
$ fastboot flash userdata && fastboot flashall
```

Monitor crash logs

```
$ adb logcat
```

Record the unique crashes

- The uniqueness is determined based on ASAN backtrace information

# Disclosure

---

Discovered vulnerabilities

# Discovered vulnerabilities – Summary

---

As of October 1, 2017, total 30 vulnerabilities have been discovered

- 13 vulnerabilities are duplicate 😞
- 17 vulnerabilities(11 critical, 5 high and 1 moderate) have been disclosed on Android Security Bulletins 😊
- Some issues are still in process 😞

Covered multiple memory corruption vulnerability types

- Heap overflow
- Heap use after free
- Stack buffer overflow
- Global buffer overflow - **Not fix yet**
- Alloc dealloc mismatch
- FPE

# Discovered vulnerabilities – Details

---

| CVE           | Type                 | Severity | components    |
|---------------|----------------------|----------|---------------|
| CVE-2017-0678 | heap-user-after-free | Critical | mp4/container |
| CVE-2017-0714 | heap-buffer-overflow | Critical | h263/codec    |
| CVE-2017-0719 | heap-buffer-overflow | Critical | mpeg2/codec   |
| CVE-2017-0718 | heap-buffer-overflow | Critical | mpeg2/codec   |
| CVE-2017-0722 | heap-buffer-overflow | Critical | h263/codec    |
| CVE-2017-0720 | heap-buffer-overflow | Critical | hevc/codec    |
| CVE-2017-0745 | heap-buffer-overflow | Critical | mpeg4/codec   |
| CVE-2017-0758 | heap-buffer-overflow | Critical | hevc/codec    |
| CVE-2017-0760 | heap-buffer-overflow | Critical | mpeg2/codec   |

# Discovered vulnerabilities – Details

---

| CVE           | Type                   | Severity | components        |
|---------------|------------------------|----------|-------------------|
| CVE-2017-0761 | heap-buffer-overflow   | Critical | avc/codec         |
| CVE-2017-0764 | stack-overflow         | Critical | vorbis/codec      |
| CVE-2017-0776 | heap-buffer-overflow   | High     | avc/codec         |
| CVE-2017-0777 | heap-buffer-overflow   | High     | sonivox/container |
| CVE-2017-0778 | heap-buffer-overflow   | High     | mp4/container     |
| CVE-2017-0820 | fpe                    | High     | mp4/container     |
| CVE-2017-0813 | alloc-dealloc-mismatch | Moderate | mp4/container     |
| CVE-2017-0814 | heap-buffer-overflow   | High     | vorbis/codec      |

# Discovered vulnerabilities – POCs

---

<https://github.com/ele7enxxh/poc-exp>

# Conclusion

---

Presentation conclusion



# Conclusion

---

A new efficient fuzzing strategy for Stagefright has been implemented, and 17 new vulnerabilities have been found

This fuzzing strategy is based on Android Nougat and AFL-2.51b, it also compatible with other version in theory, but the details may need to be changed

Is there still have 0days in Stagefright ? - Yes, but need more powerful technologies or tricks

Not only fuzzing Stagefright on Linux, Fuzzing all Android native binaries on Linux too

# Reference

---

<https://source.android.com/security/bulletin/>

<http://events.linuxfoundation.org/sites/events/files/slides/ABS2015.pdf>

<https://www.blackhat.com/docs/us-15/materials/us-15-Drake-Stagefright-Scary-Code-In-The-Heart-Of-Android.pdf>

<https://hitcon.org/2016/CMT/slide/day2-r2-c-1.pdf>

<http://lcamtuf.coredump.cx/AFL/>

[https://github.com/huntcve/slides/blob/master/seven\\_shen\\_shakacon.pdf](https://github.com/huntcve/slides/blob/master/seven_shen_shakacon.pdf)

<https://www.blackhat.com/docs/eu-16/materials/eu-16-Jurczyk-Effective-File-Format-Fuzzing-Thoughts-Techniques-And-Results.pdf>

# Thank you very much

---

Zinuo Han

ele7enxxh@gmail.com